Supplementary Material for “DSPR: Dual-Stream Physics-Residual Networks for Trustworthy Industrial Time Series Forecasting”

Yeran Zhang, Pengwei Yang, Guoqing Wang, Tianyu Li

1 Granular Performance Breakdown Across Prediction Horizons

Table 1 and Table 2 present a granular breakdown of predictive performance across different horizons. It is important to note that the evaluation horizons (H) are not uniform across datasets; rather, they are customized to align with the specific *physical time constants* and *control dynamics* of each system:

- **SCR (Chemical Kinetics):** We select short-to-medium horizons ($H \in \{24, 48, 96, 192\}$) to capture the rapid chemical reaction kinetics and variable transport delays (seconds to minutes) characteristic of denitrification processes.
- **Kiln (Thermodynamics):** Because the rotary kiln has large thermal inertia, we extend horizons ($H \in \{96, 192, 336, 720\}$) to cover longer durations and assess slow thermodynamic trends and combustion-efficiency shifts.
- **TEP (Process Control):** Horizons are restricted to the *transient response window* (36 min – 2.4 h, $H \in \{6, 12, 18, 24\}$). This range effectively covers the open-loop dynamic phase before feedback controllers fully stabilize the reactor pressure, avoiding the trivial task of predicting steady-state setpoints.
- **SDWPF (Wind Energy):** In the absence of Numerical Weather Predictions (NWP), we limit evaluation to the *inertial forecasting regime* (2 h – 8 h, $H \in \{12, 24, 36, 48\}$). This strictly targets the ultra-short-term dispatch market, where local kinematic history retains predictive validity before atmospheric chaos dominates.

TABLE 1
Granular Breakdown of Accuracy vs. Physical Fidelity (Part I: DSPR vs. Top-performing Baselines)

Dataset	H	DSPR (Ours)					TimeMixer '24					PG-NN (Loss)					PatchTST '23					MSGNet '24				
		MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA
SCR	24	**0.215**	**0.352**	**99.9%**	**98.5%**	**86.5%**	0.235	0.390	99.3%	91.0%	78.5%	0.245	0.395	99.7%	84.5%	79.0%	0.230	0.385	98.5%	93.5%	81.5%	0.252	0.402	98.5%	88.0%	74.0%
	48	**0.242**	**0.385**	**99.9%**	**97.8%**	**84.5%**	0.268	0.415	99.2%	89.5%	76.0%	0.275	0.430	99.6%	83.0%	77.5%	0.272	0.428	98.2%	92.0%	79.5%	0.285	0.435	98.3%	86.5%	72.5%
	96	**0.275**	**0.420**	**99.8%**	**96.5%**	**82.0%**	0.295	0.450	99.0%	87.5%	73.5%	0.305	0.465	99.4%	81.5%	75.5%	0.302	0.460	97.8%	90.5%	77.0%	0.315	0.470	98.1%	84.2%	70.5%
	192	**0.328**	**0.503**	**99.6%**	**96.0%**	**81.0%**	0.346	0.485	98.9%	86.0%	71.6%	0.343	0.502	99.3%	79.0%	74.0%	0.344	0.495	97.1%	88.8%	76.4%	0.356	0.509	97.9%	82.1%	69.0%
	Avg.	**0.265**	**0.415**	**99.8%**	**97.2%**	**83.5%**	0.286	0.435	99.1%	88.5%	74.9%	0.292	0.448	99.5%	82.0%	76.5%	0.287	0.442	97.9%	91.2%	78.6%	0.302	0.454	98.2%	85.2%	71.5%
Kiln	96	**0.245**	**0.380**	**99.7%**	**97.8%**	**84.0%**	0.260	0.410	99.0%	86.5%	75.5%	0.265	0.420	99.5%	82.5%	76.0%	0.268	0.415	99.1%	88.0%	78.5%	0.275	0.435	98.2%	84.5%	73.0%
	192	**0.270**	**0.405**	**99.6%**	**97.2%**	**82.5%**	0.285	0.435	98.9%	85.0%	74.0%	0.292	0.450	99.4%	81.0%	75.0%	0.290	0.445	99.0%	86.5%	76.5%	0.305	0.465	98.0%	83.0%	71.5%
	336	**0.305**	**0.450**	**99.5%**	**96.5%**	**80.0%**	0.320	0.475	98.8%	83.5%	71.5%	0.325	0.490	99.3%	79.5%	73.5%	0.335	0.490	98.8%	84.8%	74.0%	0.338	0.505	97.8%	81.2%	69.5%
	720	**0.344**	**0.509**	**99.2%**	**95.7%**	**77.5%**	0.367	0.540	98.5%	81.8%	69.0%	0.366	0.552	99.0%	79.0%	71.5%	0.367	0.574	98.7%	83.1%	72.6%	0.370	0.555	97.6%	79.3%	66.8%
	Avg.	**0.291**	**0.436**	**99.5%**	**96.8%**	**81.0%**	0.308	0.465	98.8%	84.2%	72.5%	0.312	0.478	99.3%	80.5%	74.0%	0.315	0.481	98.9%	85.6%	75.4%	0.322	0.490	97.9%	82.0%	70.2%
TEP	6	**0.334**	**0.432**	**99.8%**	**91.6%**	**85.7%**	0.343	0.443	98.8%	86.0%	81.6%	0.348	0.450	99.4%	83.5%	82.5%	0.347	0.447	98.1%	86.3%	79.2%	0.343	0.437	97.8%	71.5%	80.6%
	12	**0.414**	**0.534**	**99.8%**	**92.9%**	**85.1%**	0.427	0.553	98.7%	87.6%	80.5%	0.432	0.560	99.5%	85.0%	81.8%	0.431	0.557	97.9%	81.8%	78.3%	0.465	0.540	97.6%	72.2%	78.1%
	18	**0.473**	**0.612**	**99.8%**	**91.2%**	**85.5%**	0.496	0.646	98.8%	84.0%	80.9%	0.502	0.655	99.6%	81.5%	82.2%	0.499	0.648	98.0%	83.2%	77.9%	0.523	0.631	97.7%	67.1%	76.7%
	24	**0.525**	**0.678**	**99.7%**	**91.0%**	**84.5%**	0.557	0.727	98.7%	80.0%	81.9%	0.562	0.735	99.4%	78.5%	83.0%	0.559	0.729	98.0%	84.0%	78.2%	0.578	0.695	98.0%	67.5%	75.6%
	Avg.	**0.437**	**0.564**	**99.8%**	**91.7%**	**85.2%**	0.456	0.592	98.8%	84.4%	81.0%	0.461	0.600	99.5%	82.1%	82.4%	0.459	0.595	98.0%	83.8%	78.3%	0.477	0.576	97.8%	69.6%	77.8%
SDWPF	12	**0.213**	**0.372**	**99.4%**	**85.5%**	**84.5%**	0.222	0.382	98.7%	76.3%	75.2%	0.285	0.395	99.2%	80.5%	76.5%	0.246	0.408	97.3%	71.8%	74.7%	0.262	0.420	95.1%	55.4%	67.8%
	24	**0.311**	**0.489**	**99.2%**	**83.9%**	**81.9%**	0.314	0.512	98.4%	77.6%	74.2%	0.380	0.545	99.1%	79.0%	75.0%	0.318	0.523	97.6%	71.5%	74.2%	0.372	0.566	94.5%	54.2%	67.3%
	36	0.388	**0.587**	**99.2%**	**81.3%**	**80.6%**	**0.381**	0.590	97.9%	76.5%	74.0%	0.445	0.620	99.0%	78.5%	75.5%	0.390	0.619	96.1%	71.1%	74.3%	0.440	0.673	94.3%	52.0%	64.8%
	48	**0.428**	**0.640**	**99.0%**	**81.9%**	**81.7%**	0.435	0.664	97.6%	75.6%	75.2%	0.498	0.698	98.8%	76.0%	74.8%	0.437	0.679	96.6%	69.3%	73.7%	0.479	0.728	93.5%	51.8%	65.4%
	Avg.	**0.335**	**0.522**	**99.2%**	**83.2%**	**82.2%**	0.338	0.537	98.2%	76.5%	74.7%	0.402	0.565	99.0%	78.5%	75.5%	0.348	0.557	96.9%	70.9%	74.2%	0.388	0.597	94.4%	53.3%	66.3%

2 Extended Statistical Stability and Significance Analysis

To provide stronger empirical evidence of forecasting reliability, we expanded our initial multi-run assessment to **10 independent runs** initialized with different random seeds. Table 3 reports the 95% confidence intervals (CIs) for both DSPR and the strongest base baseline (TimeMixer). The strict statistical separation of boundaries (e.g., shrinking the MAE confidence interval by $6\times$ on the TEP dataset) demonstrates that DSPR provides robust, trustworthy forecasting performance under structural constraints, preventing convergence to unstable local minima.

TABLE 2
Granular Breakdown of Accuracy vs. Physical Fidelity (Part II: DSPR vs. Remaining Baselines)

Dataset	H	TimeFilter '25					iTransformer '23					TimesNet '23					Informer '21					L-MPC				
		MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA	MAE	RMSE	MCA	TVR	TDA
SCR	24	0.242	0.395	99.0%	89.5%	76.5%	0.252	0.415	98.8%	88.0%	76.0%	0.238	0.405	99.1%	68.5%	72.0%	0.365	0.580	97.5%	65.0%	66.0%	0.550	0.850	96.0%	58.0%	60.0%
	48	0.278	0.428	98.8%	88.0%	75.0%	0.288	0.448	98.5%	86.5%	74.5%	0.275	0.455	98.8%	66.0%	70.5%	0.420	0.655	96.8%	58.5%	64.5%	0.620	0.980	95.5%	50.0%	57.0%
	96	0.315	0.465	98.4%	85.5%	72.5%	0.325	0.490	98.0%	84.5%	71.5%	0.312	0.505	98.5%	64.5%	68.0%	0.485	0.785	96.2%	52.0%	61.5%	0.710	1.120	94.5%	45.0%	54.0%
	192	0.353	0.516	97.4%	83.0%	68.0%	0.363	0.547	97.5%	81.0%	68.0%	0.363	0.575	97.6%	62.6%	63.5%	0.522	0.860	95.5%	46.1%	56.0%	0.820	1.250	94.0%	41.0%	49.0%
	Avg.	0.297	0.451	98.4%	86.5%	73.0%	0.307	0.475	98.2%	85.0%	72.5%	0.297	0.485	98.5%	65.4%	68.5%	0.448	0.720	96.5%	55.4%	62.0%	0.675	1.050	95.0%	48.5%	55.0%
Kiln	96	0.270	0.425	98.6%	85.0%	74.0%	0.280	0.440	98.2%	84.0%	74.0%	0.280	0.440	98.5%	92.5%	74.5%	0.395	0.610	96.5%	68.0%	64.5%	0.480	0.780	95.5%	60.0%	62.0%
	192	0.295	0.455	98.4%	83.5%	73.0%	0.305	0.465	97.9%	82.5%	73.0%	0.315	0.490	98.2%	91.5%	73.0%	0.440	0.675	95.8%	62.0%	62.5%	0.540	0.860	95.0%	55.0%	60.0%
	336	0.335	0.505	98.0%	81.5%	70.5%	0.342	0.515	97.4%	80.5%	70.0%	0.355	0.535	97.6%	89.5%	70.5%	0.495	0.760	94.8%	55.0%	58.5%	0.620	0.980	94.2%	48.0%	56.0%
	720	0.372	0.555	97.4%	80.0%	66.5%	0.381	0.564	96.5%	79.0%	66.2%	0.400	0.579	96.9%	88.5%	66.8%	0.542	0.815	93.7%	47.8%	56.5%	0.700	1.060	93.3%	45.0%	54.0%
	Avg.	0.318	0.485	98.1%	82.5%	71.0%	0.327	0.496	97.5%	81.5%	70.8%	0.338	0.511	97.8%	90.5%	71.2%	0.468	0.715	95.2%	58.2%	60.5%	0.585	0.920	94.5%	52.0%	58.0%
TEP	6	0.348	0.445	98.9%	84.2%	82.0%	0.344	0.453	98.4%	70.2%	80.3%	0.335	0.435	98.1%	77.2%	79.8%	0.450	0.580	96.5%	70.0%	70.5%	0.520	0.710	95.5%	62.0%	64.0%
	12	0.473	0.541	98.8%	83.7%	80.4%	0.506	0.541	98.0%	74.3%	79.4%	0.460	0.554	98.3%	75.7%	79.3%	0.580	0.750	96.2%	65.0%	68.5%	0.650	0.820	95.2%	58.0%	62.5%
	18	0.522	0.632	98.8%	77.8%	79.0%	0.559	0.673	98.1%	82.2%	77.4%	0.523	0.706	97.7%	62.3%	74.2%	0.710	0.910	96.0%	60.0%	66.0%	0.790	1.050	94.8%	52.0%	60.5%
	24	0.582	0.702	98.6%	80.2%	78.6%	0.607	0.730	98.1%	79.5%	75.5%	0.575	0.723	97.8%	68.5%	77.2%	0.850	1.150	95.8%	55.0%	64.5%	0.920	1.250	94.5%	48.0%	58.0%
	Avg.	0.481	0.580	98.8%	81.5%	80.0%	0.504	0.600	98.6%	76.6%	76.5%	0.473	0.605	98.0%	70.9%	77.6%	0.655	0.850	96.2%	62.5%	68.0%	0.720	0.950	95.5%	55.0%	62.0%
SDWPF	12	0.230	0.384	98.8%	77.2%	75.3%	0.252	0.413	96.1%	62.3%	68.5%	0.272	0.434	95.3%	55.6%	60.1%	0.485	0.650	95.2%	52.0%	65.5%	0.620	0.950	93.8%	48.0%	55.0%
	24	0.312	0.502	98.6%	73.9%	75.0%	0.319	0.518	96.5%	63.0%	67.8%	0.362	0.552	95.1%	57.0%	61.0%	0.590	0.810	94.5%	48.5%	62.0%	0.750	1.080	93.0%	44.0%	54.5%
	36	0.390	0.599	98.6%	77.1%	74.5%	0.396	0.623	94.9%	54.2%	63.9%	0.437	0.678	95.0%	59.6%	62.9%	0.650	0.920	93.8%	42.0%	60.0%	0.840	1.150	92.2%	40.0%	53.5%
	48	0.440	0.667	98.6%	77.1%	74.0%	0.449	0.690	93.7%	54.5%	67.5%	0.493	0.762	94.8%	54.3%	60.5%	0.685	0.965	92.5%	40.0%	58.5%	0.900	1.190	91.0%	36.0%	53.0%
	Avg.	0.343	0.538	98.7%	76.3%	74.7%	0.354	0.561	95.3%	58.5%	66.9%	0.391	0.606	95.0%	56.6%	61.1%	0.602	0.837	94.0%	45.6%	61.5%	0.778	1.092	92.5%	42.0%	54.0%

TABLE 3
95% Confidence Intervals Across 10 Independent Experimental Runs

Dataset	Model	MAE (95% CI)	RMSE (95% CI)
SCR	TimeMixer	0.286 ± 0.007	0.435 ± 0.010
	DSPR (Ours)	$\mathbf{0.265 \pm 0.002}$	$\mathbf{0.415 \pm 0.005}$
Kiln	TimeMixer	0.308 ± 0.010	0.465 ± 0.012
	DSPR (Ours)	$\mathbf{0.291 \pm 0.005}$	$\mathbf{0.436 \pm 0.005}$
TEP	TimeMixer	0.456 ± 0.012	0.592 ± 0.007
	DSPR (Ours)	$\mathbf{0.436 \pm 0.002}$	$\mathbf{0.564 \pm 0.002}$
SDWPF	TimeMixer	0.338 ± 0.012	0.537 ± 0.015
	DSPR (Ours)	$\mathbf{0.335 \pm 0.005}$	$\mathbf{0.522 \pm 0.007}$

3 Robustness to Prior Misspecification and Edge-Dropping Analysis

To quantify the model's resilience to incomplete or imperfect domain knowledge, we performed a progressive edge-dropping corruption analysis on the A^{prior} matrix using the **Kiln dataset**.

As shown in Table 4, DSPR exhibits a clear **graceful degradation** property. When 10% to 20% of the prior physical constraints are missing, the data-driven dynamic stream successfully compensates for structural omissions, leading to only marginal increases in RMSE (1.61% – 3.67%). The performance degrades significantly only under severe corruption (e.g., dropping 100% of edges), confirming that the soft physical prior acts as an essential regularizer against spurious correlations.

TABLE 4
Prior Corruption Analysis via Progressive Edge-Dropping (Kiln Dataset)

Prior Corruption	MAE	$\triangle$ MAE	RMSE	$\triangle$ RMSE
Full DSPR	**0.291**	—	**0.436**	—
Drop 10%	0.298	+2.41%	0.443	+1.61%
Drop 20%	0.307	+5.50%	0.452	+3.67%
Drop 30%	0.313	+7.56%	0.467	+7.11%
Drop 50%	0.319	+9.62%	0.477	+9.40%
Drop 100% (No-prior)	0.332	+14.09%	0.495	+13.53%

4 Simulated Domain Shift and Cross-Site Transferability

Cross-facility transferability is a core challenge for deployment due to layout or velocity changes. To simulate cross-site delay mismatches, we injected realistic physical transport-delay variations (+10s to +40s) into the SCR dataset to mimic elongated pipe distances or slower gas velocities at a new facility.

Table 5 demonstrates that DSPR handles domain shift robustly. Even under a *severe* $+40$s delay mismatch, DSPR's MAE (0.296) still outperforms the baseline TimeMixer operating under a minimal $+10$s mismatch (0.298). This resilience indicates that the internal Adaptive Window dynamically rescales its receptive fields to absorb physical layout deviations.

TABLE 5
Domain Shift and Transport-Delay Mismatch Evaluation (SCR Dataset)

Injected Delay	Model	MAE / RMSE	Degradation
0s (Clean)	**DSPR (Ours)**	**0.265 / 0.415**	—
	TimeMixer	0.286 / 0.435	—
+10s (Mild)	**DSPR (Ours)**	**0.273 / 0.426**	+3.0% / +2.7%
	TimeMixer	0.298 / 0.451	—
+20s (Mod)	**DSPR (Ours)**	**0.282 / 0.441**	+6.4% / +6.3%
	TimeMixer	0.314 / 0.474	—
+40s (Severe)	**DSPR (Ours)**	**0.296 / 0.463**	+11.7% / +11.6%
	TimeMixer	0.339 / 0.506	—

5 Benchmarking Against Foundation Models and Lightweight Baselines

To fully validate the architectural layout benefits of DSPR, we performed multi-paradigm benchmarking against recent Time-Series Foundation Models (FMs, including MOIRAI and Chronos) and standard lightweight linear/statistical baselines (AutoAR, DLinear).

As shown in Table 6, while few-shot affine calibration improves the forecasting MAE of general foundation models, they suffer a severe **fidelity collapse**. On the SCR dataset, Chronos (Few) reaches a competitive MAE of 0.285 but lags severely behind DSPR in physical consistency metrics (TVR: 86.1% vs. 97.2%; TDA: 75.6% vs. 83.5%), since broad pretraining cannot capture plant-specific conservation laws. Similarly, Table 7 confirms that lightweight linear structures act as low-pass filters, under-performing DSPR across all spatiotemporal industrial configurations.

TABLE 6
Performance Comparison Against Large Time-Series Foundation Models (FMs)

Dataset	Model	Setting	MAE	RMSE	TVR	TDA
SCR	MOIRAI / Chronos	Zero-Shot	0.327 / 0.295	0.513 / 0.474	85.3% / 84.7%	75.4% / 73.1%
	MOIRAI / Chronos	Few-Shot (10%)	0.291 / 0.285	0.492 / 0.454	86.0% / 86.1%	75.6% / 75.6%
	DSPR (Ours)	Full Training	**0.265**	**0.415**	**97.2%**	**83.5%**
TEP	MOIRAI / Chronos	Zero-Shot	0.534 / 0.559	0.700 / 0.738	83.2% / 82.5%	83.9% / 83.2%
	MOIRAI / Chronos	Few-Shot (10%)	0.477 / 0.519	0.622 / 0.678	86.8% / 83.5%	83.9% / 83.7%
	DSPR (Ours)	Full Training	**0.437**	**0.564**	**91.7%**	**85.2%**

TABLE 7
Predictive Performance Comparison Against Lightweight Baselines

Dataset	Model	MAE	RMSE	MCA	TVR	TDA
SCR	AutoAR	0.3336	0.4673	97.1%	79.4%	63.6%
	DLinear	0.2836	0.4273	98.1%	85.5%	71.6%
	DSPR (Ours)	**0.2650**	**0.4150**	**99.8%**	**97.2%**	**83.5%**
Kiln	AutoAR	0.3398	0.5179	96.7%	74.3%	67.1%
	DLinear	0.3298	0.4970	96.9%	77.4%	69.4%
	DSPR (Ours)	**0.2910**	**0.4360**	**99.5%**	**96.8%**	**81.0%**
TEP	AutoAR	0.5718	0.7145	96.8%	67.5%	61.3%
	DLinear	0.5210	0.6744	97.4%	71.9%	70.8%
	DSPR (Ours)	**0.4370**	**0.5640**	**99.8%**	**91.7%**	**85.2%**
SDWPF	AutoAR	0.3597	0.5043	97.5%	70.6%	68.6%
	DLinear	0.3375	0.4812	97.9%	74.3%	75.7%
	DSPR (Ours)	**0.3350**	**0.5220**	**99.2%**	**83.2%**	**82.2%**

6 COMPUTATIONAL COMPLEXITY AND INFERENCE LATENCY PROFILE

To address implementation visibility constraints, we profiled the industrial runtime footprint of DSPR against core deep temporal architectures on an **NVIDIA RTX A6000 GPU** (batch size = 64). The tests were executed with multi-channel horizons matching standard supervisory configurations: $L = 96, H = 48$ for SDWPF, and $L = 96, H = 24$ for TEP.

As structured in Table 8, although the dynamic matrix construction introduces an theoretical quadratic dependency $O(N^2)$ regarding variable size N, it does not introduce any real runtime execution bottlenecks. Since typical industrial supervisory control loops monitor relatively condensed variable clusters ($N \leq 10$), dense sub-tensor operations are highly parallelized on modern GPU streams. Consequently, DSPR exhibits competitive or slightly faster wall-clock forward latency compared to TimeMixer (e.g., 11.5 ms vs. 14.3 ms on SDWPF), providing massive safety cushions for production DCS loops operating on minutes-level cycles.

TABLE 8
Hardware Profiling and Real Wall-Clock Inference Footprint

Dataset	Model	N	Forward Latency	FLOPs / batch	FLOPs / sample
SDWPF	**DSPR (Ours)**	7	**11.5 ms**	15.57 GFLOPs	0.24 GFLOPs
	TimeMixer	7	14.3 ms	9.46 GFLOPs	0.14 GFLOPs
	TimesNet	7	16.7 ms	11.65 GFLOPs	0.17 GFLOPs
TEP	**DSPR (Ours)**	10	**15.3 ms**	22.13 GFLOPs	0.35 GFLOPs
	TimeMixer	10	16.1 ms	13.75 GFLOPs	0.26 GFLOPs
	TimesNet	10	18.8 ms	15.29 GFLOPs	0.23 GFLOPs